

Day 1: Session 2

Data Wrangling and Visualization

Session Objectives

1. Get Started with H.E.L.P. Dataset
 2. Selecting and Exporting Data
 - A. Selecting Exporting Data (variables and subsets)
 - B. Exporting Data (variables and subsets)
 3. Getting descriptive statistics
 4. Getting started with visualization
-

0. Prework - Before You Begin

- R Packages Needed for this session:
 - `tidyverse` packages
 - * `readr`
 - * `dplyr`
 - * `ggplot2`
 - * `haven`
 - `ggtext`
- R Code Files
 - `module_d1_02.R`
- SAS Code Files
 - `sascode_d1_02.sas`
- R Data Files

- [help.Rdata](#)
 - SAS Data Files
 - [help.sas7bdat](#)
 - Other Data Files - OPTIONAL
 - [help.csv](#)
 - [helpmkh.sav](#)
-

1. Get Started with H.E.L.P. Dataset

For this course we will be working with the **HELP** dataset from this BOOK:

- “SAS and R: Data Management, Statistical Analysis, and Graphics (second edition)” by Ken Kleinman and Nicholas J. Horton
- [HELP Dataset](#)
- This data was collected for the H.E.L.P. (Health Evaluation and Linkage to Primary Care) trial, (Samet et al. 2003)

Options for importing/loading the HELP Dataset


The easiest option is to simply load the data in the ***.Rdata** native binary format, which is provided at the [HELP dataset link](#).

```
# OPTION 1: Load HELP dataset in .Rdata format  
load("help.RData")
```


The dataset is also provided in the ***.sas7bdat** SAS data format. To load SAS formatted dataset, we can use the **read_sas()** function from the [haven](#) package.

```
# OPTION 2: Load HELP dataset in .sas7bdat format  
# Click File/Import Dataset/From SAS  
library(haven)  
helpmkh <- read_sas("help.sas7bdat", NULL)
```

Another way to import this SAS formatted dataset is to use the menu functions “File/Import Dataset/From SAS”:



This opens a window that defaults to files located in the current project folder. From here we can select the `help.sas7bdat` file. We could go ahead and click the “Import” button OR we can copy the R code generated in the Code Preview window, which is the code we have above.




We can also use the `haven` package to use the `read_sav()` function to read in a SPSS formatted dataset as well.

```
# OTHER FORMATS: Load HELP dataset from SPSS
# I edited this file to add the "codebook"
# variable and value labels included
library(haven)
helpmkh2 <- read_sav("helpmkh.sav")
```


The CSV (comma separated value) file format is also readily available for many datasets as well. We can use the `read_csv()` function from the `readr` package.

```
# Read in CSV File - "Import Dataset/From Text (readr)"
library(readr)
help <- read_csv("help.csv")
```

After loading the dataset (`help.Rdata`), the dataset should now be visible in the Global Environment window at the top right:



We can take a look at what is in this dataset, by clicking on the little blue circle with an arrow in it:



This provides the same info as running `str()` function. Notice that we can see there are 453 observations (rows) and 88 variables (columns). We can also see that variables like `id`, `e2b1` are “int” (integers) and `pcs1` and `mcs1` are “num” (numeric) type variables. We can also see that `helpdata` is a `data.frame` object.

We can also get more information on this dataset using these functions.

```
# find out what kind of object helpdata is
class(helpdata)
```

```
[1] "data.frame"
```

```
# what are the dimension, rows -x- cols
dim(helpdata)
```

```
[1] 453  88
```

```
# what are all of the names of the variables
names(helpdata)
```

```
[1] "id"          "e2b1"        "g1b1"        "i11"         "pcs1"
[6] "mcs1"        "cesd1"       "indtot1"     "drugrisk1"   "sexrisk1"
[11] "pcrec1"      "e2b2"        "g1b2"        "i12"         "pcs2"
[16] "mcs2"        "cesd2"       "indtot2"     "drugrisk2"   "sexrisk2"
[21] "pcrec2"      "e2b3"        "g1b3"        "i13"         "pcs3"
[26] "mcs3"        "cesd3"       "indtot3"     "drugrisk3"   "sexrisk3"
[31] "pcrec3"      "e2b4"        "g1b4"        "i14"         "pcs4"
[36] "mcs4"        "cesd4"       "indtot4"     "drugrisk4"   "sexrisk4"
[41] "pcrec4"      "a15a"        "a15b"        "d1"          "e2b"
[46] "f1a"         "f1b"         "f1c"         "f1d"         "f1e"
[51] "f1f"         "f1g"         "f1h"         "f1i"         "f1j"
[56] "f1k"         "f1l"         "f1m"         "f1n"         "f1o"
[61] "f1p"         "f1q"         "f1r"         "f1s"         "f1t"
[66] "g1b"         "i1"          "i2"          "age"         "treat"
[71] "homeless"    "pcs"         "mcs"         "cesd"         "indtot"
[76] "pss_fr"      "drugrisk"    "sexrisk"     "satreat"
"drinkstatus"
[81] "daysdrink"  "anysubstatus" "daysanysub" "linkstatus"  "dayslink"
[86] "female"      "substance"    "racegrp"
```

💡 Get variable names with and without quotes

There are times when you need quotes and times when you do not need quotes. The default `names()` function adds the quotes. If you would like to remove them you can use the `noquote()` function.

```
# put the output from names()
# into noquote() to remove the quotes
noquote(names(helpdata))
```

```
[1] id          e2b1          g1b1          i11           pcs1
[6] mcs1        cesd1         indtot1       drugrisk1     sexrisk1
[11] pcrec1      e2b2         g1b2         i12           pcs2
[16] mcs2        cesd2         indtot2       drugrisk2     sexrisk2
[21] pcrec2      e2b3         g1b3         i13           pcs3
[26] mcs3        cesd3         indtot3       drugrisk3     sexrisk3
[31] pcrec3      e2b4         g1b4         i14           pcs4
[36] mcs4        cesd4         indtot4       drugrisk4     sexrisk4
[41] pcrec4      a15a         a15b         d1            e2b
[46] f1a         f1b          f1c          f1d           f1e
[51] f1f         f1g          f1h          f1i           f1j
[56] f1k         f1l          f1m          f1n           f1o
[61] f1p         f1q          f1r          f1s           f1t
[66] g1b         i1           i2           age           treat
[71] homeless    pcs          mcs          cesd          indtot
[76] pss_fr      drugrisk     sexrisk      satreat       drinkstatus
[81] daysdrink   anysubstatus daysanysub   linkstatus    dayslink
[86] female      substance    racegrp
```

We can also “view” the dataset, by clicking on the little table icon at the far right. Clicking the icon runs the View() function which loads the viewer window.


The screenshot shows the RStudio interface with the following components:

- Console:** Contains the command `> View(helpdata)` with a red arrow pointing to it.
- Environment:** Shows the `helpdata` object with 453 observations and 88 variables. A red arrow points to the table icon on the right side of the Environment pane.
- Files:** Shows a list of files in the `R4SASUsers` directory, including `dat2.sas7bdat`, `help.csv`, `help.Rdata`, `help.sas7bdat`, `helpfmt_ageswomen.sas7bdat`, and `helpfmt_race.sas7bdat`.

SAS Environment

Remember that we started with creating an RStudio Project for the folder `C:\R4SASUsers`. This process is similar to beginning your SAS compute session by setting a `LIBNAME` statement, so that SAS knows where your files are for your current project.

And I typically use the `DATA` step to make a copy of my data into the `WORK` library for the current computation session. The `WORK` library is like the `Global Environment` in R.



2. Selecting and Exporting Data

Now that we have the HELP dataset loaded into our computing session, we can start working with the `helpdata` dataset object.

2A. Selecting Data (variables and subsets)

Selecting Variables or Columns - Base R approach

Using base R we can “select” variables out of the `helpdata` dataset using the `$` dollar sign.

```
# select the age variable  
# from helpdata and show in Console  
helpdata$age
```

```
[1] 37 37 26 39 32 47 49 28 50 39 34 58 58 60 36 28 35 29 27 27 41 33 34 31  
48  
[26] 34 35 34 29 35 43 37 29 33 20 38 28 33 40 43 28 45 42 30 34 36 44 41 30  
37  
[51] 35 37 44 47 38 37 34 41 29 35 36 27 36 40 38 42 26 41 43 28 30 42 22 31  
30  
[76] 25 26 35 53 29 32 24 35 32 47 26 45 33 45 33 27 40 40 37 26 27 35 29 33  
39  
[101] 33 35 38 44 28 33 30 35 32 42 37 41 28 30 35 35 41 37 30 39 32 50 33 27  
33  
[126] 38 43 24 35 49 49 33 24 28 45 46 37 32 45 39 34 32 32 31 45 30 36 25 48  
42  
[151] 33 36 41 30 57 57 47 54 55 33 29 33 28 37 28 32 31 36 39 29 38 33 31 39  
33  
[176] 31 46 36 22 33 35 38 28 33 49 43 33 29 34 41 47 24 31 40 32 32 39 19 49  
27  
[201] 38 32 22 36 32 35 35 41 36 43 45 39 47 32 33 39 44 35 31 25 48 35 42 51  
55  
[226] 32 41 34 30 34 38 41 31 29 36 45 36 30 40 27 39 39 37 43 20 35 32 42 27  
30  
[251] 27 41 32 47 36 32 33 30 29 34 34 40 45 37 32 26 31 39 49 45 43 38 23 35  
23  
[276] 42 29 43 29 39 32 35 22 39 38 56 36 40 22 39 47 32 41 32 37 41 31 33 30  
32  
[301] 35 32 33 30 44 46 43 47 34 47 40 34 48 37 35 38 27 39 23 35 53 31 32 33  
25  
[326] 37 26 29 30 47 33 36 23 36 34 28 33 26 30 41 31 28 59 39 36 47 26 22 36  
34
```

```
[351] 27 34 21 33 42 46 26 36 47 48 32 38 43 30 40 38 22 39 22 37 37 44 38 37
43
[376] 39 45 39 31 32 42 33 47 24 27 38 53 39 32 27 43 31 41 27 28 39 39 21 29
31
[401] 29 45 25 24 41 27 21 27 31 41 33 49 41 25 41 34 29 28 29 36 36 24 38 31
26
[426] 35 26 33 46 33 39 27 33 36 23 33 26 38 52 39 36 44 37 33 31 25 31 24 33
49
[451] 39 59 45
```

If we don't want to simply see the result in the "Console" window, we can save the result into a new object to use later.

```
# save as separate object
ages <- helpdata$age
```

dplyr approach using pipes

A better way is to use the **dplyr** package which uses a stepwise computing approach with the **magrittr** pipe `%>%`. What is a "pipe"?

In computer programming a "pipe" allows you to take the output from the preceding command and "send it" as input to the next command.

Let's take a look at the help page for the **select()** function from the **dplyr** package.

Use **select(helpdata, age)** to select the **age** variable from the **helpdata** dataset.

```
# tidyverse - dplyr Approach using select()
library(dplyr)

# using the dplyr::select() function
# the first argument is the dataset
# which is defined as .data = helpdata
# select the "age" variable from helpdata
select(helpdata, age)
```

```
   age
1   37
2   37
3   26
4   39
5   32
6   47
```

7	49
8	28
9	50
10	39
11	34
12	58
13	58
14	60
15	36
16	28
17	35
18	29
19	27
20	27
21	41
22	33
23	34
24	31
25	48
26	34
27	35
28	34
29	29
30	35
31	43
32	37
33	29
34	33
35	20
36	38
37	28
38	33
39	40
40	43
41	28
42	45
43	42
44	30
45	34
46	36
47	44
48	41
49	30

50	37
51	35
52	37
53	44
54	47
55	38
56	37
57	34
58	41
59	29
60	35
61	36
62	27
63	36
64	40
65	38
66	42
67	26
68	41
69	43
70	28
71	30
72	42
73	22
74	31
75	30
76	25
77	26
78	35
79	53
80	29
81	32
82	24
83	35
84	32
85	47
86	26
87	45
88	33
89	45
90	33
91	27
92	40

93	40
94	37
95	26
96	27
97	35
98	29
99	33
100	39
101	33
102	35
103	38
104	44
105	28
106	33
107	30
108	35
109	32
110	42
111	37
112	41
113	28
114	30
115	35
116	35
117	41
118	37
119	30
120	39
121	32
122	50
123	33
124	27
125	33
126	38
127	43
128	24
129	35
130	49
131	49
132	33
133	24
134	28
135	45

136 46
137 37
138 32
139 45
140 39
141 34
142 32
143 32
144 31
145 45
146 30
147 36
148 25
149 48
150 42
151 33
152 36
153 41
154 30
155 57
156 57
157 47
158 54
159 55
160 33
161 29
162 33
163 28
164 37
165 28
166 32
167 31
168 36
169 39
170 29
171 38
172 33
173 31
174 39
175 33
176 31
177 46
178 36

179 22
180 33
181 35
182 38
183 28
184 33
185 49
186 43
187 33
188 29
189 34
190 41
191 47
192 24
193 31
194 40
195 32
196 32
197 39
198 19
199 49
200 27
201 38
202 32
203 22
204 36
205 32
206 35
207 35
208 41
209 36
210 43
211 45
212 39
213 47
214 32
215 33
216 39
217 44
218 35
219 31
220 25
221 48

222 35
223 42
224 51
225 55
226 32
227 41
228 34
229 30
230 34
231 38
232 41
233 31
234 29
235 36
236 45
237 36
238 30
239 40
240 27
241 39
242 39
243 37
244 43
245 20
246 35
247 32
248 42
249 27
250 30
251 27
252 41
253 32
254 47
255 36
256 32
257 33
258 30
259 29
260 34
261 34
262 40
263 45
264 37

265 32
266 26
267 31
268 39
269 49
270 45
271 43
272 38
273 23
274 35
275 23
276 42
277 29
278 43
279 29
280 39
281 32
282 35
283 22
284 39
285 38
286 56
287 36
288 40
289 22
290 39
291 47
292 32
293 41
294 32
295 37
296 41
297 31
298 33
299 30
300 32
301 35
302 32
303 33
304 30
305 44
306 46
307 43

308 47
309 34
310 47
311 40
312 34
313 48
314 37
315 35
316 38
317 27
318 39
319 23
320 35
321 53
322 31
323 32
324 33
325 25
326 37
327 26
328 29
329 30
330 47
331 33
332 36
333 23
334 36
335 34
336 28
337 33
338 26
339 30
340 41
341 31
342 28
343 59
344 39
345 36
346 47
347 26
348 22
349 36
350 34

351 27
352 34
353 21
354 33
355 42
356 46
357 26
358 36
359 47
360 48
361 32
362 38
363 43
364 30
365 40
366 38
367 22
368 39
369 22
370 37
371 37
372 44
373 38
374 37
375 43
376 39
377 45
378 39
379 31
380 32
381 42
382 33
383 47
384 24
385 27
386 38
387 53
388 39
389 32
390 27
391 43
392 31
393 41

394 27
395 28
396 39
397 39
398 21
399 29
400 31
401 29
402 45
403 25
404 24
405 41
406 27
407 21
408 27
409 31
410 41
411 33
412 49
413 41
414 25
415 41
416 34
417 29
418 28
419 29
420 36
421 36
422 24
423 38
424 31
425 26
426 35
427 26
428 33
429 46
430 33
431 39
432 27
433 33
434 36
435 23
436 33

```
437 26
438 38
439 52
440 39
441 36
442 44
443 37
444 33
445 31
446 25
447 31
448 24
449 33
450 49
451 39
452 59
453 45
```

However, we can use a pipe programming approach and re-write this line of code in 2 steps by listing the dataset first “and then” selecting the “age” variable from this dataset. Notice the `%>%` pipe operator at the end of the first line of code.

```
helpdata %>%
  select(age)
```

The following 2 figures show that the `%>%` pipe operator takes the output `helpdata` and “sends it” into the `select()` function as the first argument. The second argument `age` which is the variable to be selected is then listed in the 2nd step of the code sequence after the `%>%` pipe operator.

R

```
arrange(helpdata, age)
```

```
helpdata %>% arrange(age)
```

**R**

```
arrange(helpdata, age)
```

```
helpdata %>% arrange(age)
```



Load data → then → arrange it by variable

We can continue using the `%>%` pipe operator to put together multiple code steps in sequence where the output from the preceding function call is used as input to the next function call.

This block of code does the following:

1. load the `helpdata` dataset “and then”
2. selects the `age` variable from `helpdata` “and then”
3. shows the top 10 rows of ages using the `help()` function.

```
helpdata %>%  
  select(age) %>%  
  head(10)      # show top 10 rows
```

	age
1	37
2	37
3	26
4	39
5	32
6	47
7	49
8	28
9	50
10	39

We can also save the ages from the HELP dataset into another object called `ages2` but this time using the `dplyr` `select()` function.


```
# save object
ages2 <- helpdata %>%
  select(age)
```

Programming pipes `%>%` vs `|>`

For the purpose of this course, I will stick with the `magrittr` pipe `%>%`. However, there is also a “base R” pipe `|>`. For the most part these are interchangeable. however, there are a few differences discussed [in this blog post](#).

Base R versus `dplyr` approaches - differences in output

Both the “base R” and `dplyr` approaches will select the ages of the participants from the dataset. However, the resulting `ages` or `ages2` objects are slightly different. We can see how they look in the Global Environment window and by comparing the results from the `class()` function.



The “base R” approach results in a vector of integer numbers. The `dplyr` approach results in a simple `data.frame` object with 1 column - the `dplyr` approach preserves the original object type of `helpdata`.

```
# base R approach result
class(ages)
```

```
[1] "integer"
```

```
# dplyr approach result
class(ages2)
```

```
[1] "data.frame"
```

We can always select more than one variable at a time by simply listing them separated by a comma. Notice that we do not need the quotes around the variable names.

```
# select multiple variables
agecesd <- helpdata %>%
  select(age, cesd)
```

Selecting Rows - Base R

Another “Base R” method to select variables includes using indexes. For example select rows 1:10 and column 69 (which is `age`). We could also specify the “`age`” column name using “base R” syntax of `data[rownum, colnum]` or `data[rownum, colname]`.

```
# Base R approach
# pull out rows 1-10 and column 69 (which is age)
# dataset[row,column] index approach
helpdata[1:10, 69]
```

```
[1] 37 37 26 39 32 47 49 28 50 39
```

```
# another option, use variable name
helpdata[1:10, "age"]
```

```
[1] 37 37 26 39 32 47 49 28 50 39
```

We could also filter out the males and only keep the rows of data for the females. The “Base R” approach gets a little tricky:

```
# select women, get stats for age
# female coded as
#   0 for men
#   1 for women
# Use == for logical check of equality
# learn more
# help('==', package = "base")
ageswomen <- helpdata[helpdata$female == 1, "age"]
```

! The difference between == and =

R uses a single equal sign = for assigning a value to an argument in a function. For example, `plot(x=xdata, y=ydata)` to assign the variable to the x and y axes in a plot. But if we want to check to see if two objects are the same or match in a logical sense, we need to use == to see if two things match or not, `helpdata$female == 1`. This results in values of TRUE or FALSE which are reserved logical values in R. R then uses the values of TRUE or FALSE to parse which rows to keep or remove.

Selecting Rows - dplyr Approach

As you’ve seen above the base R approach gets tricky quickly. This is why learning the `dplyr` approach is good - the code is much easier to understand.

There are 2 main approaches for choosing a subset of rows out of a dataset using `dplyr` - either the `slice()` or `filter()` functions.

In the output below, again notice the difference between the “Base R” and `dplyr` approach output for `ageswomen` and `ageswomen2`.

```
# dplyr approach =====
# notice the output is different again
# get rows 1-10 and column 69 (which is age)
helpdata %>%
  slice(1:10) %>%
  select(69)
```

```
   age
1   37
2   37
```

3	26
4	39
5	32
6	47
7	49
8	28
9	50
10	39

```
# another option, use variable name
helpdata %>%
  slice(1:10) %>%
  select(age)
```

```
      age
1    37
2    37
3    26
4    39
5    32
6    47
7    49
8    28
9    50
10   39
```

```
# select women, get ages
ageswomen2 <- helpdata %>%
  filter(female == 1) %>%
  select(age)

class(ageswomen)
```

```
[1] "integer"
```

```
class(ageswomen2)
```

```
[1] "data.frame"
```

In summary, the `dplyr` programming approach using pipes `%>%` is easier to read and understand than the more complicated nested syntax of the Base R approach. The `dplyr` webpage has a nice article summarizing these differences at <https://dplyr.tidyverse.org/articles/base.html#one-table-verbs>.

SAS Code for quick comparison

Here is the SAS code that also does the same data selection examples as above.

- `firstobs` and `obs` options are used instead of the `dplyr slice()` function above
- `keep` is similar to the `dplyr select()` function
- for the logical filter step, we can use the SAS `where` statement

```
data helpfmt_rows1_10;
  set helpfmt(firstobs=1 obs=10); /* keep rows 1-10*/
  keep age cesd; /* Only keep age and cesd */
run;


* save ages for women;
data helpfmt_ageswomen;
  set helpfmt;
  where female = 1; /* keep women */
  keep age; /* Only keep age */
run;
```

2B. Exporting Data (variables and subsets)

Now that we have created two new datasets, `agecesd` and `ageswomen2`, we could save these as R binary `*.Rdata` files for use in future computing sessions.

```
# now that we have created some new datasets (data.frame)
# let's save agecesd and ageswomen2 as .Rdata files
save(agecesd,
      file = "agecesd.Rdata")
save(ageswomen2,
      file = "ageswomen2.Rdata")
```

Check your local file folder for this project - you should now have two new datafiles:



You could also save more than one object at a time into an R binary datafile - this can be very useful. These objects do NOT have to be datasets - they could be ANY kind of R object (models, graphic plots, custom functions and more).

```
save(ages, ages2, ageswomen,
      file = "otherobjects.Rdata")
```

Let's remove these objects and reload them. We can use the list `ls()` function to see all of the object currently loaded into the Global Environment for this R computation session. The `rm()` function removes the objects we specify.

```
# let's remove these objects and reload them
# list objects in Global Environment
ls()
```

```
[1] "agecesd"    "ages"      "ages2"     "ageswomen" "ageswomen2"
[6] "help"      "helpdata"  "helpmkh"   "helpmkh2"
```

```
rm(ages, ages2, ageswomen, agecesd, ageswomen2)
ls()
```

```
[1] "help"      "helpdata" "helpmkh"   "helpmkh2"
```

Now load back `agecesd` from `agecesd.Rdata` and `ages`, `ages2`, `ageswomen` from `otherobjects.Rdata`.

```
# load "agecesd.Rdata" a single data object
load("agecesd.Rdata")
ls()
```

```
[1] "agecesd"  "help"     "helpdata" "helpmkh"  "helpmkh2"
```

```
# load "otherobjects.Rdata", multiple objects
load("otherobjects.Rdata")
ls()
```

```
[1] "agecesd"  "ages"     "ages2"     "ageswomen" "help"      "helpdata"
[7] "helpmkh"  "helpmkh2"
```

Suppose you want to keep everything (all of the objects) in your current Global Environment, you can also use `save.image()` which will save everything so you don't have to list every object individually.

```
# if you want to save everything in your  
# global environment at the moment  
save.image(file = "allobjects.Rdata")
```

Remove everything and then load it all back.

```
# remove everything  
rm(list = ls())  
ls()
```

```
character(0)
```

```
# load it back  
load("allobjects.Rdata")  
ls()
```

```
[1] "agecesd"  "ages"      "ages2"      "ageswomen" "help"      "helpdata"  
[7] "helpmkh"  "helpmkh2"
```

SAS Code for quick comparison

```
* write these new datasets to your project folder ;

data rforsas.helpfmt_rows1_10;
    set work.helpfmt_rows1_10;
run;

data rforsas.helpfmt_ageswomen;
    set work.helpfmt_ageswomen;
run;

* Look at rforsas Library;

* look at WORK Library;

proc datasets lib=work nolist;
    delete helpfmt_rows1_10 helpfmt_ageswomen;
quit;

* read saved datasets back in to
  WORK library if wanted;

data work.helpfmt_ageswomen;
    set rforsas.helpfmt_ageswomen;
run;
```

3. Getting descriptive statistics

R - Quick `summary()` statistics

Besides the `class()` function for figuring out what an object's type is, the `summary()` function is another workhorse in the R language.

Generic functions

Both `class()` and `summary()` are what are called generic functions. A generic function “... invokes particular methods which depend on the class of the first argument.” We will discuss this further tomorrow when we use the `summary()` function after running a regression model.

For now, let's use the `summary()` function on a variable from a dataset. Technically this runs `summary.default()`, see `help(summary.default)`.

Let's get some summary statistics for `age` from `helpdata`.

```
# Get descriptive statistics with summary()
# get min, max, mean, median and quartiles
summary(helpdata$age)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
19.00	30.00	35.00	35.65	40.00	60.00

This is a really helpful quick way to get summary statistics. I just wish the original programmers had also included the standard deviation in the output which would give me the parametric statistics: mean and sd as well as the non-parametric stats: median, 25th percentile and 75th percentile along with the min and max.

Here is the `dplyr` approach:

```
helpdata %>%
  select(age) %>%
  summary()
```

age	
Min.	:19.00
1st Qu.	:30.00
Median	:35.00
Mean	:35.65
3rd Qu.	:40.00
Max.	:60.00

SAS - summary statistics with PROC UNIVARIATE

There are many different ways to get summary descriptive statistics for variables in SAS. I often use PROC UNIVARIATE for this purpose.

```
* get descriptive statistics of age;
proc univariate data=help plots;
  var age;
run;
```

R - More individual functions for descriptive stats

Let's also get the standard deviation using `sd()`

```
# base R
sd(helpdata$age)
```

```
[1] 7.710266
```

```
# dplyr pipe approach
helpdata %>%
  select(age) %>%
  sd()
```

```
Error in `is.data.frame()` :
! 'list' object cannot be coerced to type 'double'
```

OOPS This throws an error. Let's figure out why.

```
# the sd() function is expecting
# a vector type variable not a data.frame
# what is the result after the select() step
```

```
aa <- helpdata %>%
  select(age)
class(aa)
```

```
[1] "data.frame"
```

See `help(sd)` to see that the `sd()` function is expecting a numeric vector as input and NOT a `data.frame`. So we need to either convert the output using `unlist()` after `select()` or use a different selection function like `pull()` from `dplyr`. *NOTE: A `data.frame` is a special type of `list` class object in R.*

```
# use the unlist() function to convert it to a vector
helpdata %>%
  select(age) %>%
  unlist() %>%
  sd()
```

```
[1] 7.710266
```

The `dplyr` package also has the `pull()` function which results in a vector and NOT a `data.frame` output object.

```
# the dplyr package also has
# the pull() function which results in
# a vector as well
helpdata %>%
  pull(age) %>%
  sd()
```

```
[1] 7.710266
```

i `select()` versus `pull()` in `dplyr` package

Take a moment and compare the output “Value” produced from the `pull()` versus `select()` functions - look at the “help” window for both functions and scroll down to the “Value” section of the “help” page.

Notice that the output from `pull()` is a vector class type and the output from `select()` is the same class type as the `.data` argument.

Output from `pull()`

Value

A vector the same size as `.data`.

Output from `select()`**Value**

An object of the same type as `.data`. The output has the following properties:

- Rows are not affected.
- Output columns are a subset of input columns, potentially with a different order. Columns will be renamed if `new_name = old_name` form is used.
- Data frame attributes are preserved.
- Groups are maintained; you can't select off grouping variables.

What if we have missing data indicated by NA in our dataset?

Let's generate a numeric vector with a missing value and then see what happens with these functions.

```
# vector of numeric values
x <- c(1,2,3,4,5,NA)
x
```

```
[1] 1 2 3 4 5 NA
```

```
# get summary statistics
summary(x)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NAs
1	2	3	3	4	5	1

This worked, but notice the summary output clearly shows we have 1 NA value.

What about the `sd()` function?

```
# try the sd() function  
sd(x)
```

```
[1] NA
```

The value returned was NA. This can be frustrating, but there is a method to the madness. Technically, since there is a value missing, we have to explicitly tell the `sd()` function what we want to do - do we want to remove the missing values and then compute the standard deviation. We can do that by setting the argument for `na.rm` to `TRUE`.

It is also good that we get NA in our output since this tells us we have missing data in our dataset for this variable which prompts us to make decisions on how to handle or adjust for the missing data BEFORE completing the analysis or computations.

```
# now tell sd what to do if there are missing values  
sd(x, na.rm=TRUE)
```

```
[1] 1.581139
```

```
# dplyr approach  
x %>%  
  sd(na.rm=TRUE)
```

```
[1] 1.581139
```

It is a good idea to just get in the habit of adding `na.rm=TRUE` to these functions to make sure you are providing explicit instructions on how to handle the missing data.

R - parametric and non-parametric stats

Base R code adding `na.rm=TRUE`:

```
# adjust for missing values  
# parametric stats  
mean(helpdata$age, na.rm=TRUE)
```

```
[1] 35.65342
```

```
sd(helpdata$age, na.rm=TRUE)
```

```
[1] 7.710266
```

```
# get non-parametric statistics  
min(helpdata$age, na.rm=TRUE)
```

```
[1] 19
```

```
max(helpdata$age, na.rm=TRUE)
```

```
[1] 60
```

```
median(helpdata$age, na.rm=TRUE)
```

```
[1] 35
```

```
quantile(helpdata$age, probs = 0.25, na.rm=TRUE)
```

```
25%  
30
```

```
quantile(helpdata$age, probs = 0.75, na.rm=TRUE)
```

```
75%  
40
```

 Default settings!

NOTE: Pay attention to default settings! For the `quantile()` function to compute percentiles, be sure to check the default settings for this function. There are 9 possible algorithms you can choose from. The default is `type=7`. This may not match other software packages default algorithms. For large sample sizes these differences should be minor. But for smaller sample sizes, these algorithms from computing percentiles may/will be different.

Let's look at the help window for `quantile()`.

```
help(quantile, package = "stats")
```

Quantile Help Windows

```
## Default S3 method:
quantile(x, probs = seq(0, 1, 0.25), na.rm = FALSE,
        names = TRUE, type = 7, digits = 7,
        fuzz = if(type == 7L) 0 else 4 * .Machine$double.eps,
        ...)
```

Arguments

x	numeric vector whose sample quantiles are wanted, or an object of a class for which a method has been defined (see also 'details'). NA and NaN values are not allowed in numeric vectors unless <code>na.rm</code> is TRUE.
probs	numeric vector of probabilities with values in $[0, 1]$. (Values up to '2e-14' outside that range are accepted and moved to the nearby endpoint.)
na.rm	logical; if true, any NA and NaN values are removed from x before the quantiles are computed.
names	logical; if true, the result has a names attribute. Set to FALSE for speedup with many probs.
→ type	an integer between <u>1 and 9</u> selecting one of the <u>nine quantile algorithms</u> detailed below to be used.
digits	used only when <code>names</code> is true: the precision to use when formatting the percentages. In R versions up to 4.0.x, this had been set to <code>max(2, getOption("digits"))</code> , internally.
fuzz	small non-negative number to protect against rounding errors when <code>j <- floor(np + m)</code> , (<code>np</code> "a version of" <code>n * probs</code> , see the formula below), is computed.

Types

`quantile` returns estimates of underlying distribution quantiles based on one or two order statistics from the supplied elements in `x` at probabilities in `probs`. One of the nine quantile algorithms discussed in [Hyndman and Fan \(1996\)](#), selected by `type`, is employed.

All sample quantiles are defined as weighted averages of consecutive order statistics. Sample quantiles of type i are defined by:

$$Q_i(p) = (1 - \gamma)x_j + \gamma x_{j+1}$$

where $1 \leq i \leq 9$, $\frac{j-m}{n} \leq p < \frac{j-m+1}{n}$, x_j is the j -th order statistic, n is the sample size, the value of γ is a function of $j = \lfloor np + m \rfloor$ and $g = np + m - j$, and m is a constant determined by the sample quantile type.

Discontinuous sample quantile types 1, 2, and 3

For types 1, 2 and 3, $Q_i(p)$ is a discontinuous function of p , with $m = 0$ when $i = 1$ and $i = 2$, and $m = -1/2$ when $i = 3$.

Type 1

Inverse of empirical distribution function. $\gamma = 0$ if $g = 0$, and 1 otherwise.

Type 2

Similar to type 1 but with averaging at discontinuities. $\gamma = 0.5$ if $g = 0$, and 1 otherwise (SAS default, see [Wicklin \(2017\)](#)).

Compare the output differences between using algorithm `type=7` to `type=9`:

```
# get 25th, 50th, 75th percentiles
# use default type=7
quantile(helpdata$age,
         probs = c(0.25, 0.50, 0.75),
         na.rm=TRUE)
```

```
25% 50% 75%
30  35  40
```

```
# compare to type=9
quantile(helpdata$age,
         probs = c(0.25, 0.50, 0.75),
         type=9,
         na.rm=TRUE)
```

```
      25%      50%      75%
30.0000 35.0000 40.3125
```

SAS - parametric and non-parametric stats

To produce neatly formatted tables of summary statistics, I also like using PROC MEANS in SAS.

```
* another approach with proc means
- get parametric statistics;
proc means data=help n min max mean std;
  var age;
run;

* another approach with proc means
- get non-parametric statistics;
proc means data=help n q1 median q3;
  var age;
run;

* There are 5 options for quantile
  calculation in SAS, see SAS help for
  "Quantile and Related Statistics"
  QNTLDEF=5 is the default
  compare to QNTLDEF=4;

proc means data=help q1 median q3 QNTLDEF=4;
  var age;
run;
```

R - dplyr approach for custom statistics

dplyr approach to show multiple custom statistics all in the same output.

```
# create custom output
helpdata %>%
  summarise(
    age_n = n(),
    meanage = mean(age, na.rm=TRUE),
    sdage = sd(age, na.rm=TRUE),
    medage = median(age, na.rm=TRUE),
    q25age = quantile(age,
                      probs = 0.25,
                      na.rm=TRUE),
    q75age = quantile(age,
                      probs = 0.75,
                      na.rm=TRUE),
    minage = min(age, na.rm=TRUE),
    maxage = max(age, na.rm=TRUE)
  )
```

	age_n	meanage	sdage	medage	q25age	q75age	minage	maxage
1	453	35.65342	7.710266	35	30	40	19	60

Let's leverage other functions in `dplyr` to help. Let's get the custom statistics for the women only by adding `filter(female == 1)`.

```
# get age stats for women only
helpdata %>%
  filter(female == 1) %>%
  summarise(
    age_n = n(),
    meanage = mean(age, na.rm=TRUE),
    sdage = sd(age, na.rm=TRUE),
    medage = median(age, na.rm=TRUE),
    q25age = quantile(age,
                      probs = 0.25,
                      na.rm=TRUE),
    q75age = quantile(age,
                      probs = 0.75,
                      na.rm=TRUE),
    minage = min(age, na.rm=TRUE),
    maxage = max(age, na.rm=TRUE)
  )
```

	age_n	meanage	sdage	medage	q25age	q75age	minage	maxage
1	107	36.25234	7.584858	35	31	40.5	21	58

We can also use the `group_by` function to get stratified output by group as well. Let's look at these statistics for each of the 4 race/ethnicity groups:

```
# get age stats by racegrp
helpdata %>%
  group_by(racegrp) %>%
  summarise(
    age_n = n(),
    meanage = mean(age, na.rm=TRUE),
    sdage = sd(age, na.rm=TRUE),
    medage = median(age, na.rm=TRUE),
    q25age = quantile(age,
                      probs = 0.25,
                      na.rm=TRUE),
    q75age = quantile(age,
                      probs = 0.75,
                      na.rm=TRUE),
    minage = min(age, na.rm=TRUE),
    maxage = max(age, na.rm=TRUE)
  )
```

```
# A tibble: 4 x 9
  racegrp age_n meanage sdage medage q25age q75age minage maxage
<fct>   <int>   <dbl> <dbl> <dbl> <dbl> <dbl> <int> <int>
1 black     211    35.7  7.08    35    31    39      20     60
2 hispanic   50    33.2  7.99    32   28.2   36.2     21     55
3 other      26    35.0  7.66    34    30   40.5     22     48
4 white     166    36.5  8.28    36    30    42     19     58
```

SAS - Custom stats tables for filtered subsets and stratified by group

To get statistics for a subset of the data, we can use the `WHERE` statement in SAS. To get stratified output, we can also use the `CLASS` statement.

```
* Get stats for a subset  
ages and cesd for women only;
```

```
proc means data=help;  
  where female = 1;  
  var age;  
run;
```

```
* Get stats for a subset  
ages and cesd BY racegrp;
```

```
proc means data=help;  
  class racegrp;  
  var age;  
run;
```

R - Sorting and displaying data

Let's combine multiple `dplyr` functions to wrangle our data to pull out specific parts of our dataset.

For example lets:

1. just look at the women
2. look at only the subject `id` and `cesd` scores
3. sort the data by `age`

```
# select IDs and cesd scores for women
# sort by age
helpdata %>%
  filter(female == 1) %>%
  select(id, cesd) %>%
  arrange(age)
```

```
Error in `arrange()` :
i In argument: `..1 = age`.
Caused by error:
! object 'age' not found
```

Oops we got an error - why didn't this work?

We selected `id` and `cesd` in step 2 so when we got ready to sort the data in step 3 we had already removed `age`. We need to change the order of these steps.

I've also added a 4th step to only show the `head()` or top rows for the 5 youngest women in the output below.

```
# select women
# sort by age
# show IDs, ages and cesd scores for 5 youngest women
helpdata %>%
  filter(female == 1) %>%
  arrange(age) %>%
  select(id, age, cesd) %>%
  head(5)
```

	id	age	cesd
1	155	21	39
2	342	22	51
3	324	23	48

```
4 141 24 45
5 171 24 34
```

What if we want to see the `id` and `cesd` scores for the 5 oldest women. We could sort the data in descending order by age OR we could keep the ascending order and look at the bottom or `tail()` part of the dataset. The results of these 2 options is slightly different, but the result is for the same 5 id's.

```
# show IDs, ages and cesd scores for 5 oldest women
# option 1
helpdata %>%
  filter(female == 1) %>%
  arrange(desc(age)) %>%
  select(id, age, cesd) %>%
  head(5)
```

```
      id age cesd
1    12  58   49
2   181  57   39
3   264  55   30
4     9  50   50
5   134  50   35
```

```
# option 2 - same results slightly different order
helpdata %>%
  filter(female == 1) %>%
  arrange(age) %>%
  select(id, age, cesd) %>%
  tail(5)
```

```
      id age cesd
103    9  50   50
104  134  50   35
105  264  55   30
106  181  57   39
107   12  58   49
```

SAS - Sorting and displaying data

Here is the SAS code to do the same sorting and displaying of the selected variables for the women in `helpdata`.

```
* sorting and displaying data =====;
* sort data by age for women;

proc sort data=help out=womensortage;
    by age;
    where female = 1;
run;

* show ids and cesd scores
  for 5 youngest women;

proc print data=womensortage(obs=5);
    var id age cesd;
run;

* show ids and cesd scores
  for 5 oldest women
  have to know number of rows=107
  set firstobs=102
  set last observation to obs=107;

proc print data=womensortage(firstobs=102 obs=107);
    var id age cesd;
run;

* or sort descending and reuse code above;

proc sort data=help out=womensortdage;
    by descending age;
    where female = 1;
run;

proc print data=womensortdage(obs=5);
    var id age cesd;
run;
```

R - Table summaries of categorical variables and associated tests

Another workhorse function in R is the `table()` function. Let's look at the breakdown of gender in this dataset:

```
# get frequencies and percents
# table of gender, show any missing values
# base R
table(helpdata$female,
      useNA = "ifany")
```

```
0    1
346 107
```

```
# dplyr approach
helpdata %>%
  select(female) %>%
  table(useNA = "ifany")
```

```
female
0    1
346 107
```

💡 Handling missing values for categorical data

Similar to using `na.rm=TRUE` for the statistical functions above, we can/should also tell `table()` what we want to do with the missing values. By adding `useNA = "ifany"` we will also get a count of how many missing values there are for this categorical variable.

```
# table of categorical data with NAs
fruit <- c("apple","apple","apple",
          "grape", NA, "grape")
table(fruit)
```

```
fruit
apple grape
3       2
```

```
table(fruit, useNA = "ifany")
```

```
fruit
apple grape <NA>
      3      2      1
```

The `e2b` variable in the `HELP` dataset comes from a branching logic question about whether or not the participant had entered detox and, if so, how many times they have been in detox. So, this variable has missing data for the people who did NOT enter detox. Let's look at the output with and without a listing of the missing NA's.

```
# table summary without a list of the missing NAs
table(helpdata$e2b)
```

```
 1  2  3  4  5  6  7  8 11 14 17 21
90 57 29 16  6  4  3  4  2  1  1  1
```

```
# table summary with a list of the missing NAs
table(helpdata$e2b, useNA = "ifany")
```

```
 1    2    3    4    5    6    7    8    11    14    17    21 <NA>
90   57   29   16    6    4    3    4    2    1    1    1   239
```

Two-way tables:

```
# table of racegrp by gender
# base R
table(helpdata$racegrp,
      helpdata$female,
      useNA = "ifany")
```

```
      0    1
black 155  56
hispanic 40  10
other   19   7
white  132  34
```

```
# dplyr
helpdata %>%
  select(racegrp, female) %>%
  table(useNA = "ifany")
```

```
      female
racegrp    0   1
  black   155 56
hispanic   40 10
  other    19  7
  white   132 34
```

R - Create a “factor” type variable with labels for gender

It can get annoying trying to remember what the various numeric codes are for each category. To combine numeric coding with value labels, we should create a “factor-type” variable. We can then use these to make easier to read and present tables for ourselves and the readers of our results.

```
# add value labels to female
# and make it a "factor"
# Base R option
helpdata$female.f <-
  factor(helpdata$female,
         levels = c(0, 1),
         labels = c("male", "female"))

# dplyr options with mutate
helpdata <- helpdata %>%
  mutate(
    gender.f = factor(female,
                     levels = c(0, 1),
                     labels = c("male", "female")))
```

Two-way table with labels added - use `gender.f`.

```
# dplyr
helpdata %>%
  select(racegrp, gender.f) %>%
  table(useNA = "ifany")
```

	gender.f	
racegrp	male	female
black	155	56
hispanic	40	10
other	19	7
white	132	34

R - A Better Formatted Table with `gmodels` package

A function I really like for making table is the `CrossTable()` function from the `gmodels` package. You can use this function to make tables that resemble the output you expect from `PROC FREQ` from SAS.

In the code shown below, I'm asking to display the:

- raw counts
- expected counts
- and column %'s
- row %, total % and chisquare proportion are turn off, set to **FALSE**
- I'm also asking to run the **chisq** test and **fisher**'s exact tests
- and the output should be formatted like SAS

Notice that I am using the `gender.f` factor type variable with better labels.

```
# make table with gmodels package
# and CrossTable() function
# show expected counts
#     percents within columns
#     chi-square test
#     fisher's exact test

library(gmodels)
gmodels::CrossTable(helpdata$racegrp,
                     helpdata$gender.f,
                     expected=TRUE,
                     prop.r=FALSE,
                     prop.t=FALSE,
                     prop.chisq=FALSE,
                     chisq=TRUE,
                     fisher=TRUE,
                     format = "SAS")
```

Cell Contents	

	N
Expected N	
N / Col Total	

Total Observations in Table: 453

	helpdata\$gender.f		
helpdata\$racegrp	male	female	Row Total
black	155	56	211
	161.161	49.839	
	0.448	0.523	
hispanic	40	10	50
	38.190	11.810	
	0.116	0.093	
other	19	7	26
	19.859	6.141	
	0.055	0.065	
white	132	34	166
	126.790	39.210	
	0.382	0.318	
Column Total	346	107	453
	0.764	0.236	

Statistics for All Table Factors

Pearson's Chi-squared test

Chi^2 = 2.423904 d.f. = 3 p = 0.4892005

Fisher's Exact Test for Count Data

Alternative hypothesis: two.sided
p = 0.4867413

Another nicely formatted table with gtsummary package

The [gtsummary package](#) provides the `tbl_cross()` function which make nicely formatted tables of summary statistics. Here is an example table for the 4 race/ethnicity groups by gender.

The steps and lines of code below do the following:

1. within the table:
 - put `racegrp` in rows
 - put `gender.f` in columns
 - only show the % within each column
2. for the overall table, make the row and column labels **bold** text
3. add the p-values for the statistical test (which defaults to the `chisquare.test()`)

```
# using gtsummary package
# and tbl_cross() function
library(gtsummary)
helpdata %>%
  tbl_cross(row = racegrp,
            col = gender.f,
            percent = "column") %>%
  bold_labels() %>%
  add_p()
```

Table 1

	gender.f		Total	p-value ¹
	male	female		
racegrp				0.5
black	155 (45%)	56 (52%)	211 (47%)	
hispanic	40 (12%)	10 (9.3%)	50 (11%)	
other	19 (5.5%)	7 (6.5%)	26 (5.7%)	
white	132 (38%)	34 (32%)	166 (37%)	
Total	346 (100%)	107 (100%)	453 (100%)	

¹Pearson's Chi-squared test

We could also get the Fisher's exact test by changing the default setting in the `add_p()` function.

```
# run fisher's exact test
helpdata %>%
  tbl_cross(row = racegrp,
            col = gender.f,
            percent = "column") %>%
  bold_labels() %>%
  add_p(test = "fisher.test")
```

Table 2

	gender.f		Total	p-value ¹
	male	female		
racegrp				0.5
black	155 (45%)	56 (52%)	211 (47%)	
hispanic	40 (12%)	10 (9.3%)	50 (11%)	
other	19 (5.5%)	7 (6.5%)	26 (5.7%)	
white	132 (38%)	34 (32%)	166 (37%)	
Total	346 (100%)	107 (100%)	453 (100%)	

¹Fisher's exact test

And here are a few more customization's that can be helpful. Changing the decimal precision.

```
# show p-value to 3 digits
# show counts to 0 decimals
# show columns percent to 2 decimal digits
helpdata %>%
  tbl_cross(row = racegrp,
            col = gender.f,
            percent = "column",
            digits = c(0, 2)) %>%
  bold_labels() %>%
  add_p(test = "fisher.test",
        pvalue_fun = label_style_pvalue(digits = 3))
```

Table 3

	gender.f		Total	p-value ¹
	male	female		
racegrp				0.487
black	155 (44.80%)	56 (52.34%)	211 (46.58%)	
hispanic	40 (11.56%)	10 (9.35%)	50 (11.04%)	
other	19 (5.49%)	7 (6.54%)	26 (5.74%)	
white	132 (38.15%)	34 (31.78%)	166 (36.64%)	
Total	346 (100.00%)	107 (100.00%)	453 (100.00%)	

¹Fisher's exact test

Adding better labels and a title for the table.

```
# add better labels
# add caption title
helpdata %>%
  tbl_cross(row = racegrp,
            col = gender.f,
            percent = "column",
            digits = c(0, 2),
            label = list(
              racegrp = "Race",
              gender.f = "Gender"
            )) %>%
  bold_labels() %>%
  add_p(test = "fisher.test",
        pvalue_fun = label_style_pvalue(digits = 3)) %>%
  modify_caption("**Table Race by Gender**")
```

Table 4: Table Race by Gender

	Gender		Total	p-value ¹
	male	female		
Race				0.487
black	155 (44.80%)	56 (52.34%)	211 (46.58%)	
hispanic	40 (11.56%)	10 (9.35%)	50 (11.04%)	
other	19 (5.49%)	7 (6.54%)	26 (5.74%)	
white	132 (38.15%)	34 (31.78%)	166 (36.64%)	
Total	346 (100.00%)	107 (100.00%)	453 (100.00%)	

¹Fisher's exact test

Back to default chi-square test results:

```
# and with chi-square test results
helpdata %>%
  tbl_cross(row = racegrp,
            col = gender.f,
            percent = "column",
            digits = c(0, 2),
            label = list(
              racegrp = "Race",
              gender.f = "Gender"
            )) %>%
  bold_labels() %>%
  add_p(pvalue_fun = label_style_pvalue(digits = 3)) %>%
  modify_caption("**Table Race by Gender**")
```

Table 5: Table Race by Gender

	Gender		Total	p-value ¹
	male	female		
Race				0.489
black	155 (44.80%)	56 (52.34%)	211 (46.58%)	
hispanic	40 (11.56%)	10 (9.35%)	50 (11.04%)	
other	19 (5.49%)	7 (6.54%)	26 (5.74%)	
white	132 (38.15%)	34 (31.78%)	166 (36.64%)	
Total	346 (100.00%)	107 (100.00%)	453 (100.00%)	

¹Pearson's Chi-squared test

SAS - Table summaries of categorical variables and associated tests

For getting summary tables for categorical data, PROC FREQ works well in SAS and also is used for performing chi-square and Fisher's exact tests.

```
* get frequency tables;

proc freq data=helpfmt;
  tables female racegrp;
run;

* get race by gender
  chi-square test
  and fisher's exact tests
  show column percents, counts,
  and expected counts;

proc freq data=helpfmt;
  table racegrp*female / chisq fisher expected norow nopercnt;
run;
```

4. Getting started with visualization

The base R `graphics` package is actually very powerful in its own right. There are some amazing “artwork” graphics you can do with R graphics.

That said, for this course we will focus on learning `ggplot2` which has been around for more than 15-20 years!!

`ggplot2` is so popular it has almost become the defacto standard for R graphics. Many hundreds of other packages have been developed that are compatible or leverage the `ggplot2` syntax and functions. As of 05/20/2026 at 4:15 pm EDT, there are 256 packages alone on [CRAN](#) starting with `gg` and there are many more R packages that do not have `gg` in the package name that are still compatible.

Handy Graphics Resources for R

Here are some good places to start to get help and get inspired for making graphics, figures and art in R.

- [Additional Resources for this Course](#)
- [Best Data Visualization Packages](#)
- [Data Visualization with Base R](#)

Basics for `ggplot2`


The `gg` stands for the “grammar of graphics”. The idea is to build a graphic using the data, layers of aesthetics, annotations, statistics, graphic objects and more.

Let’s build a scatterplot of the `mcs` with `cesd` from the `helpdata` dataset. We’ll color the points by gender and add some informative labels, title and caption as well. We’ll also add “best fit” linear regression models for the association between the `mcs` and `cesd` by gender.

1. Initialize `ggplot()` for the dataset
 - This opens a graphics window.

```
library(ggplot2)


# initialize ggplot() for the dataset
ggplot(data = helpdata)
```



2. Define the aesthetics for the x and y axes.

- this now adds axes scaled for each variable on each axis and lays out a grid.

```
ggplot(data = helpdata,  
  aes(x = mcs,  
    y = cesd))
```




3. How to “see” the data?

- We need to specify a “geom” which is a geometric object. Since we’re doing a scatterplot, we will use `geom_point()` to show the data using points or dots.

i Adding Layers with +


In `ggplot2` we build a graphical figure adding layers of “geoms” and other attributes like labels using + plus operator.

```
# scatterplot of mcs and cesd
ggplot(data = helpdata,
       aes(x = mcs,
           y = cesd)) +
  geom_point()
```



4. Add color for the points using `gender.f`

```
# color by gender
ggplot(data = helpdata,
       aes(x = mcs,
           y = cesd,
           color = gender.f)) +
  geom_point()
```



5. Add better labels using `labs()` A little text wrangling


In the long caption I added below, I inserted `\n` to add a line break. Learn more at:

- <https://r-graph-gallery.com/289-control-ggplot2-title.html>
- <https://ggplot2-book.org/annotations.html>

 Aligning aesthetic with legend title

Since we used `gender.f` for the aesthetic for `color`, when we want to change the label for the legend title we use `color` to assign the custom legend for that aesthetic - otherwise the legend title defaults to the variable name “gender.f”.

```
# add title, footnote, and better axis labels
# color by gender
# update legend title
# \n inserts line break in long caption
ggplot(data = helpdata,
       aes(x = mcs,
           y = cesd,
           color = gender.f)) +
geom_point() +
labs(
  title = "CESD Scores by MCS Scores",
  x = "CESD Scores",
  y = "MCS Scores",
  color = "Gender",
  caption = "CESD = Center for Epidemiological Studies-Depression \nand MCS
    ↪ = Mental Component Scale for SF36 Quality of Life"
)
```



6. Add “best fit” linear regression model lines to the plot by gender by adding the layer `geom_smooth(method = "lm")`.

```
# add best fit lines by gender
ggplot(data = helpdata,
       aes(x = mcs,
           y = cesd,
           color = gender.f)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(
    title = "CESD Scores by MCS Scores",
    subtitle = "Models by Gender",
    x = "CESD Scores",
    y = "MCS Scores",
    color = "Gender",
    caption = "CESD = Center for Epidemiological Studies-Depression \nand MCS
               ↪ = Mental Component Scale for SF36 Quality of Life"
  )
```



SAS - Visualizations with PROC SGPLOT

PROC SGPLOT is also really useful for making graphics visualizations in SAS. Here is the SAS code to make plots similar to what we did above.

```
* scatterplot of mcs and cesd;
proc sgplot data=helpfmt;
  scatter x=mcs y=cesd;
run;

* color by gender;
proc sgplot data=helpfmt;
  scatter x=mcs y=cesd / group=female;
run;

* add best fit lines by group;
proc sgplot data=helpfmt;
  reg x=mcs y=cesd / group=female;
run;

* add title, footnotes,
  update xaxis and yaxis labels
  and legend title
  add confidence bands for best fit lines;

TITLE1 "CESD Scores by MCS Scores";
FOOTNOTE1 "CESD = Center for Epidemiological Studies-Depression";
FOOTNOTE2 "MCS = Mental Component Scale for SF36 Quality of Life";

proc sgplot data=helpfmt;
  reg x=mcs y=cesd / clm group=female;
  xaxis label="MCS Scores";
  yaxis label="CESD Scores";
  keylegend / title="Gender";
run;

TITLE; FOOTNOTE; /* Clears the titles and footnotes */
```

OPTIONAL PLOT EXAMPLE

For the plot example below, we need to first gather some summary statistics and save them into a new dataset object. We will build the plot below using the original dataset `ToothGrowth` (which is a built-in dataset) AND the `tg.summary` data object.

We will be building this error bar plot showing the means $\pm$ standard deviations of tooth lengths over 3 dose levels of Vitamin C from either Vitamin C (VC) supplements or orange juice(OJ). And we're going to also show the data points overlaid on top of the errorbars.

We will build the plot as follows:

1. data source = `ToothGrowth`
2. aesthetics `x=dose` and `y=len`
3. to show connections between the means use `geom_line()`
4. to make the points for the means use `geom_point()`
5. to make the errorbars (+/-) standard deviations use `geom_errorbar()`
6. to overlay the original individual data points use `geom_jitter()` - within this geom I have added aesthetics to customize `color`, `shape` and `fill` for the points according to the supplement type (VC vs OJ)
7. I have also used the `shape_xxxx_manual()` functions to set the specific
 - point “shapes”, see `help(par, package = "graphics")` and scroll through the options for “pch” also see `help(points, package = "graphics")`.
 - outline “color” for which I used two of the built-in colors, see a full list by running `colors()`
 - point “fill” colors for which I used HEXCOLOR codes, you can pretty much use any hexadecimal code for color, see example codes at <https://htmlcolorcodes.com/>
8. I also added a `facet_wrap()` to create a 2 panel plot for each supplement
9. I added custom labels
 - most of these we've used in the above examples
 - for the caption I actually included some “markdown” syntax to allow for more custom text formatting - to use these formats I had to first load the `ggtext` package, and add the `theme(plot.caption = element_markdown())`
 - to customize the “legend” title, I set `color=`, `shape=` and `fill=` to the same text since all of these aesthetics are aligned with the same variable `supp` for the supplement type.


```
# OPTIONAL PLOT EXAMPLE
# first get summary stats
# build error bar using stats
# means +/- SD for error bars
#
# use layers and multiple data sources
# to make lines plus points
# and add custom shapes
# custom color outlines and fills
# by groups
#
# add ggtext package to enable markdown
# for formatting text in the caption
# see theme(plot.caption = ) below
library(ggtext)

# use ToothGrowth builtin dataset
# get summary stats
tg.summary <- ToothGrowth %>%
  group_by(dose) %>%
  summarise(
    sd = sd(len, na.rm = TRUE),
    len = mean(len)
  )
tg.summary
```

```
# A tibble: 3 x 3
  dose    sd  len
<dbl> <dbl> <dbl>
1  0.5  4.50  10.6
2    1   4.42  19.7
3    2   3.77  26.1
```

```
# Combine with jitter points
ggplot(ToothGrowth,
       aes(x=dose,
           y=len)) +
  geom_line(data = tg.summary,
            aes(group = 1),
            linewidth = 1.5) +
  geom_point(data = tg.summary,
             size = 2) +
  geom_errorbar(aes(ymin = len-sd, ymax = len+sd),
                width = 0.2,
                linewidth = 1.5,
                data = tg.summary) +
  geom_jitter(position = position_jitter(0.2),
              aes(color = supp,
                  shape = supp,
                  fill = supp),
              size = 3) +
  scale_shape_manual(values = c(21, 24)) +
  scale_color_manual(values = c("black",
                                "blue")) +
  scale_fill_manual(values = c("#cf34eb",
                                "#34eb4c")) +
  facet_wrap(vars(supp)) +
  labs(x = "Dose (mg/day)",
       y = "Length",
       title = "Guinea Pig Incisor Odontoblast Growth by Vitamin C Dose",
       subtitle = "Means +/- SD",
       caption = "Citation: Crampton, E. W. (1947). The growth of the
↪ odontoblast of the incisor teeth <br>as a criterion of vitamin C
↪ intake of the guinea pig. <br>*The Journal of Nutrition*,
↪ **33**_(5), 491-504. <br>*doi:10.1093/jn/33.5.491*.",
       color = "Supplement Type",
       shape = "Supplement Type",
       fill = "Supplement Type") +
  theme(plot.caption = element_markdown())
```

Guinea Pig Incisor Odontoblast Growth by Vitamin C Dose

Means $\pm$ SD

ation: Crampton, E. W. (1947). The growth of the odontoblast of the incisor teeth
as a criterion of vitamin C intake of the guinea pig.
The Journal of Nutrition, **33**(5), 491–504.
[doi:10.1093/jn/33.5.491](https://doi.org/10.1093/jn/33.5.491).

R and SAS Code For This Module

- [module_d1_02.R](#)
- [sascode_d1_02.sas](#)

References

- Boettiger, Carl. 2021. *Knitcitations: Citations for Knitr Markdown Files*. <https://github.com/cboettig/knitcitations>.
- Iannone, Richard, Joe Cheng, Barret Schloerke, et al. 2026. *Gt: Easily Create Presentation-Ready Display Tables*. <https://gt.rstudio.com>.
- R Core Team. 2026. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. <https://doi.org/10.32614/R.manuals>.
- Samet, Jeffrey H., Mary Jo Larson, Nicholas J. Horton, Kathleen Doyle, Michael Winter, and Richard Saitz. 2003. “Linking Alcohol- and Drug-Dependent Adults to Primary Medical Care: A Randomized Controlled Trial of a Multi-Disciplinary Health Intervention in a Detoxification Unit.” *Addiction* 98 (4): 509–16. <https://doi.org/10.1046/j.1360-0443.2003.00328.x>.
- Sjoberg, Daniel D., Joseph Larmarange, Michael Curry, et al. 2025. *Gtsummary: Presentation-Ready Data Summary and Analytic Result Tables*. <https://github.com/ddsjoberg/gtsummary>.
- Sjoberg, Daniel D., Karissa Whiting, Michael Curry, Jessica A. Lavery, and Joseph Larmarange. 2021. “Reproducible Summary Tables with the Gtsummary Package.” *The R Journal* 13: 570–80. <https://doi.org/10.32614/RJ-2021-053>.
- Warnes, Gregory R., Ben Bolker, Thomas Lumley, et al. 2022. *Gmodels: Various r Programming Tools for Model Fitting*. <https://doi.org/10.32614/CRAN.package.gmodels>.
- Wickham, Hadley. 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Winston Chang, Lionel Henry, et al. 2025. *Ggplot2: Create Elegant Data Visualisations Using the Grammar of Graphics*. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Romain François, Lionel Henry, Kirill Müller, and Davis Vaughan. 2026. *Dplyr: A Grammar of Data Manipulation*. <https://dplyr.tidyverse.org>.

Wickham, Hadley, Jim Hester, and Jennifer Bryan. 2026. *Readr: Read Rectangular Text Data*. <https://readr.tidyverse.org>.

Wickham, Hadley, Evan Miller, and Danny Smith. 2025. *Haven: Import and Export SPSS, Stata and SAS Files*. <https://haven.tidyverse.org>.

Wilke, Claus O., and Brenton M. Wiernik. 2022. *Ggtext: Improved Text Rendering Support for Ggplot2*. <https://wilkelab.org/ggtext/>.

Other Helpful Resources

[Other Helpful Resources](#)